

Linear regression

November 11, 2016

Agenda


- 1 Model representation
- 2 Cost function
- 3 Gradient descent

Useful resources

- 1 Coursera. Machine learning (Andrew Ng)
- 2 HSE course. Week 2. Week 4.

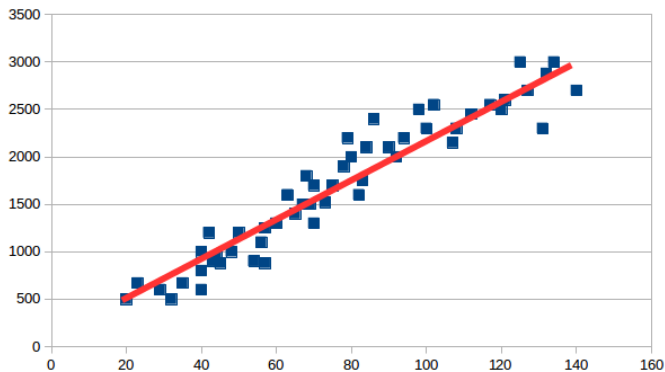
Model representation

House prices




Model representation

House prices



Model representation

Supervised learning. We give right answer for each example of data.
Regression problem - predict real-valued output.



Model representation

Training set of housing prices.

Size (x)	Price (y)
20	500
40	800
65	1300
...	...


Notation:

m = number of training example

x = input variable features


y = output variable target variable

Model representation



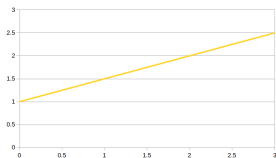
Model representation

$h_{\theta}(x) = \Theta_0 + \Theta_1 x$ Linear regression with one variable.

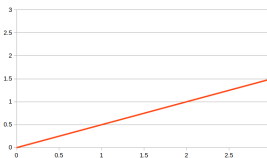


Model representation

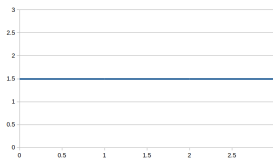
How to choose Θ_i
 $h_{\theta}(x) = \Theta_0 + \Theta_1 x$



(a) $\Theta_0 = 1$ $\Theta_1 = 0.5$



(b) $\Theta_0 = 0$ $\Theta_1 = 0.5$



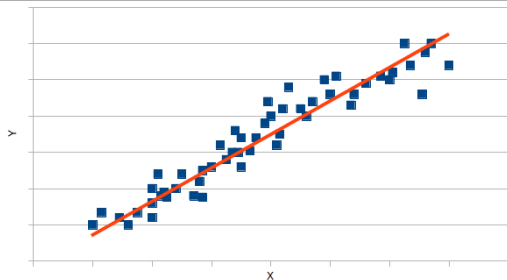
(c) $\Theta_0 = 1.5$ $\Theta_1 = 0$

Cost function

$$h_{\theta}(x) = \Theta_0 + \Theta_1 x$$


We need to choose Θ_0, Θ_1 that $h_0(x)$ is close to y for our training examples (x, y) .

So we should minimize $J(\Theta_0, \Theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_0^i(x) - y^i)^2$
 $J(\Theta_0, \Theta_1)$ - **Loss function**, squared error function.




Cost function

$h_{\theta}(x) = \Theta_0 + \Theta_1 x$ - hypothesis function.
 $\Theta_0 = 0, \Theta_1 = 0.5$



Cost function

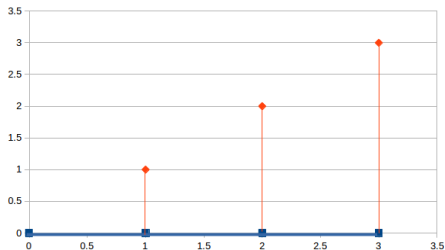
$h_{\theta}(x) = \Theta_0 + \Theta_1 x$ - hypothesis function.
 $\Theta_0 = 0, \Theta_1 = 0.5$



$$J(0, 0.5) = \frac{1}{2 \cdot 3} ((0.5 - 1)^2 + (1 - 2)^2 + (1.5 - 3)^2) = 0.58$$

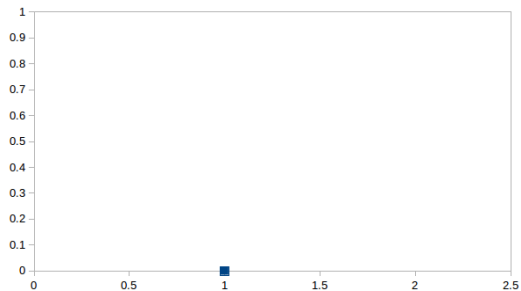
Cost function

$h_{\theta}(x) = \Theta_0 + \Theta_1 x$ - hypothesis function.
 $\Theta_0 = 0, \Theta_1 = 0$




$$J(0,0) = \frac{1}{2 \cdot 3} ((0-1)^2 + (0-2)^2 + (0-3)^2) = 2.33$$

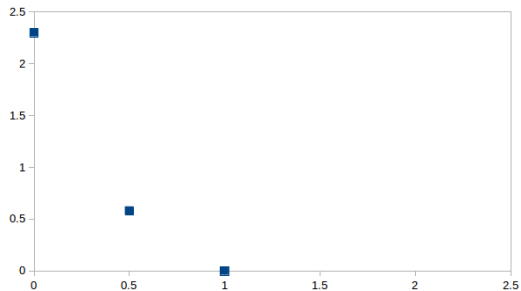
Cost function




Cost function




Cost function



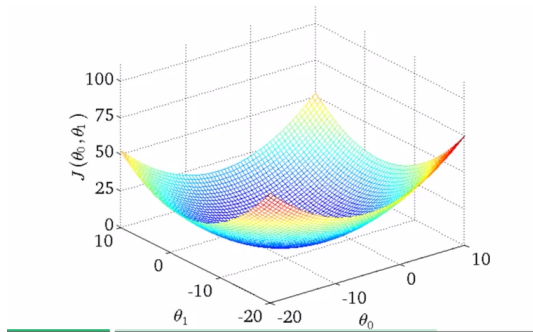
Cost function



Cost function



Cost function



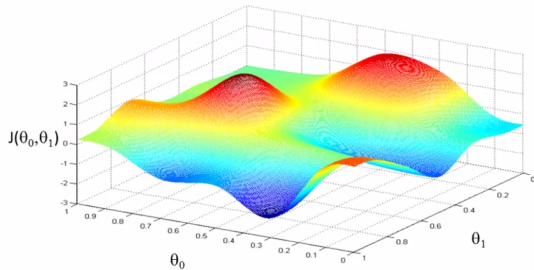
Gradient descent

Have some function $J(\Theta_0, \Theta_1)$

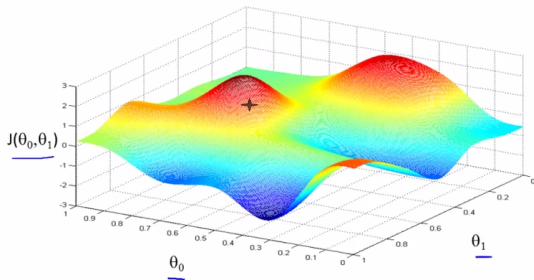
Want to find $\min J(\Theta_0, \Theta_1)$

- 1 Start from initial Θ_0, Θ_1
- 2 Keep changing Θ_0, Θ_1 to reduce $J(\Theta_0, \Theta_1)$ until we find minimum


Gradient descent




Gradient descent



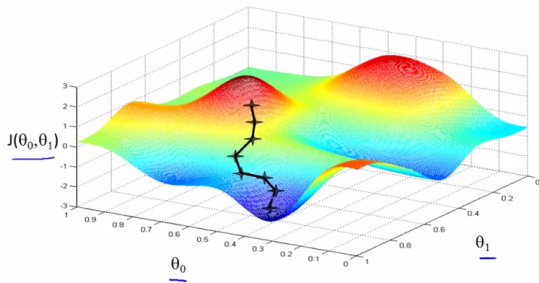
Gradient descent



Gradient descent



Gradient descent



Gradient descent

Gradient descent algorithm.

Repeat until convergence

$$\Theta_j = \Theta_j - \alpha \frac{\partial}{\partial \Theta_j} J(\Theta_0, \Theta_1)$$

Simultaneous update:

$$\text{temp0} = \text{temp0} - \alpha \frac{\partial}{\partial \Theta_0} J(\Theta_0, \Theta_1)$$


$$\text{temp1} = \text{temp1} - \alpha \frac{\partial}{\partial \Theta_1} J(\Theta_0, \Theta_1)$$

$$\Theta_0 = \text{temp0}$$


$$\Theta_1 = \text{temp1}$$

Gradient descent


$$\Theta_1 = \Theta_0 - \alpha \frac{\partial}{\partial \Theta_1} J(\Theta_0, \Theta_1)$$




Gradient descent



Gradient descent

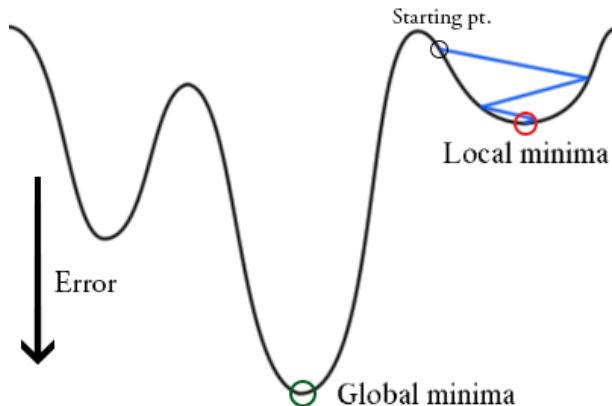


Large learning rate: Overshooting.



Small learning rate: Many iterations until convergence and trapping in local minima.

Gradient descent



Gradient descent

Gradient descent for linear regression.

Repeat until convergence

$$\Theta_j = \Theta_j - \alpha \frac{\partial}{\partial \Theta_j} J(\Theta_0, \Theta_1) = \Theta_j - \alpha \frac{\partial}{\partial \Theta_j} \frac{1}{2m} \sum_{m=1}^i (\Theta_0 + \Theta_1 x^i - y^i)^2$$

$$\Theta_0 = \Theta_0 - \alpha \frac{\partial}{\partial \Theta_j} \frac{1}{m} \sum_{m=1}^i (h_{\Theta}(x^i) - y^i)$$

$$\Theta_1 = \Theta_1 - \alpha \frac{\partial}{\partial \Theta_j} \frac{1}{m} \sum_{m=1}^i (h_{\Theta}(x^i) - y^i) x^i$$

Multiple features

Training set of housing prices.

Size (x)	Price (y)
20	500
40	800
65	1300
...	...

Notation:

m = number of training example

x = input variable features

y = output variable target variable

$$h_{\Theta}(x) = \Theta_0 + \Theta_1 x$$

Multiple features

Training set of housing prices.

Size (x)	Bedrooms	Floors	Age	Price (y)
x_1	x_2	x_3	x_4	y
20	5	1	45	500
40	3	2	20	800
65	3	3	14	1300
...

Notation:

n = number of features

x^i = input (features) of i^{th} training example.

x_j^i = value of feature j in i^{th} training example

$$h_{\Theta}(x) = \Theta_0 + \Theta_1 x + \Theta_2 x + \Theta_3 x + \Theta_4 x$$

Multiple features

$h_{\Theta}(x) = \Theta_0 + \Theta_1 x + \Theta_2 x + \Theta_3 x + \Theta_4 x$ - hypothesis function.

Let's define $x_0 = 1$, so in matrix notation:

$$h_{\Theta}(x) = \Theta^T x$$

Multiple features

Gradient descent

Hypothesis: $h_{\Theta}(x) = \Theta_0 + \Theta_1x + \Theta_2x + \Theta_3x + \dots + \Theta_nx$

Parameters: Θ — n -dimensional vector

Cost function: $J(\Theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\Theta}(x^i) - y^i)^2$

Gradient descent: Repeat $\{\Theta_j = \Theta_j - \alpha \frac{\partial}{\partial \Theta_j} J(\Theta)\}$ (simultaneously update for every $j = 0, \dots, n$)

Feature scaling


Make sure that your features are on a similar scale.

E.g. $x_1 - \text{size } (0 - 150m^2) \rightarrow x_1 = \frac{\text{size}}{150}$

$x_2 - \text{number of bedrooms } (1 - 5) \rightarrow x_2 = \frac{\text{number}}{5}$


Mean normalization

Replace x_i with $x_i - \mu_i$ to make features have approximately zero mean




Gradient descent

Making sure gradient descent is working correctly.




Polynomial regression

$$h_{\Theta}(x) = \Theta_0 + \Theta_1 x + \Theta_2 x^2 + \Theta_3 x^3 = \Theta_0 + \Theta_1(\text{size}) + \Theta_2(\text{size})^2 + \Theta_3(\text{size})^3$$



Overfitting

Overfitting refers to a model that models the training data too well. Overfitting happens when a model learns the detail and noise in the training data to the extent that it negatively impacts the performance on the model on new data.



Homework

Logistic regression

- 1 Coursera. Andrew Ng course. Week 3.
- 2 Scikit-learn documentation